

Unsupervised Learning and Evolutionary Computation Using R

Winter Term 2024/2025

Exercise Sheet 1 (21st October, 2024)

Installing the Jupyter notebook R kernel: please following [this tutorial](#)¹. Essentially, once Jupyter and R are installed executing the following commands in an interactive R session should be enough to get it running:

```
install.packages('IRkernel')  
IRkernel::installspec() # to register the kernel in the current R installation
```

Exercise 1 (Coin tossing simulation)

The function `sample(...)` can be used to sample random numbers uniformly at random from a given set. Read the documentation of the function (`?sample`). Write a simulation study on tossing a single fair coin n times for different values of $n \rightarrow \infty$. Report the absolute/relative number of head/tail in some way (functions `barplot` or `table` could be useful).

Exercise 2 (Trimmed mean)

The arithmetic mean is likely the most popular and frequently used measure of mass. A related measure is the so-called α -trimmed arithmetic mean which for a series of data points x_1, \dots, x_n and $\alpha \in [0, 1)$ is defined as

$$\bar{x}_t(x_1, \dots, x_n) = \frac{1}{n - 2\lfloor \alpha n \rfloor} \sum_{i=\lfloor \alpha n \rfloor}^{n-\lfloor \alpha n \rfloor} x_{(i)}.$$

Here, $\lfloor x \rfloor$ rounds its argument x to the nearest integer value lower than x and $x_{(i)}$ for $i = 1, \dots, n$ is the i th order statistic (the i th order statistic is the value that comes at the i th position if the data is sorted in increasing order). Hence, the trimmed mean cuts off a fraction of $\alpha\%$ lowest and highest observations and calculates the mean of the remaining ones. Implement a function `trimmedMean(x, alpha)` that implements the trimmed mean following the formula given above. Why should it be useful to use the trimmed mean instead of the standard arithmetic mean?

Exercise 3 (Re-scaling)

We will see later in the semester that for many multi-variate methods, i.e., methods that deal with data of at least two (numeric) variables $X_1, \dots, X_p, p \geq 2$, different variable scales are problematic. Here, it is often useful to transform these variables to a common scale; oftentimes $[0, 1]$. For realisations x_1, \dots, x_n of a numeric variable X this *re-scaling* can be realised by calculating

$$\tilde{x}_i = \frac{x_i - \min_j x_j}{\max_j x_j - \min_j x_j}, i = 1, \dots, n.$$

¹URL: <https://github.com/IRkernel/IRkernel>

1. Implement a function `rescale(x)` that expects a numeric vector `x` as input and returns the re-scaled version.
2. Test your function on different input vectors.
3. Imagine the input vector contains at least one so-called NA-value (NA for not available, a special value type in R). Check your input on a vector with at least one NA value, e.g., `c(1, 5, NA, 10, 3)`. Modify your function by adding a logical parameter `na.rm` which defaults to `FALSE`. If set to `TRUE` NA values should be ignored during re-scaling.
4. Imagine now the input vector potentially contains either `Inf` or `Inf`. Again, check how your implementation behaves and come up with a possible solution.

Hint: `is.infinite()` might be useful.

Exercise 4 (R loop performance)

In the presentation slides you learned that loops should be avoided in R unless there is a dependency. For a numeric vector `x` the build-in function `cumsum(x)` builds the cumulative sum vector of the same length, i.e., the i th entry of `cumsum(x)` is $\sum_{j=1}^i x_j$.

1. Write a function `cumsum_loop(x)` that calculates the cumulative sum in R using basic R loops. Why can't we use one of the `*apply` functions to solve the task?
2. Generate random numeric vectors of different length $n \in \{1\,000, 100\,000, 1\,000\,000, \dots\}$. Run both the build-in version and your implementation several times and store the elapsed running time. Visualise the running time difference, e.g., with box-plots.