

# A GENTLE INTRODUCTION INTO THE STATISTICAL PROGRAMMING LANGUAGE R


## LECTURE: UNSUPERVISED LEARNING AND EVOLUTIONARY COMPUTATION USING R

**Jakob Bossek**

MALEO Group, Department of Computer Science, Paderborn University, Germany

26<sup>th</sup> Oct, 2024

## Why are there so many slides?



Forest reaching the ocean (for the first time 😊); Forest Gump, ©Paramount pictures 1994.

*"That day, for no particular reason, I decided to go for a little run. So I ran to the end of the road. And when I got there, I thought maybe I'd run to the end of town. And when I got there, I thought maybe I'd just run across Greenbow County. And I figured, since I run this far, maybe I'd just run across the great state of Alabama. And that's what I did. I ran clear across Alabama. For no particular reason I just kept on going. I ran clear to the ocean. And when I got there, I figured, since I'd gone this far, I might as well turn around, just keep on going. When I got to another ocean, I figured, since I'd gone this far, I might as well just turn back, keep right on going."* - **Forrest Gump**

# Table of Contents

1. Introduction
2. Data types and classes
3. Output
4. Basic R graphics
5. Control flow
6. Functions
7. Advanced: Parallelization
8. Literature Recommendations
9. Wrap-Up

# Introduction

## Preface

- ▶ My original goal: at most 30 slides to teach the very essentials
- ▶ However, presentation slides evolved to a deck of  $> 100$  slides
- ▶ Consequence: we will skip a lot of stuff ☺
- ▶ Benefit for you: quite a lot important aspects are covered. Take it as a reference
- ▶ Many exercises without solutions (discuss with fellow students).<sup>1</sup>
- ▶ Enjoy! ☺



---

<sup>1</sup> 2<sup>nd</sup> exercise sheet will also contain many R exercises.

## Introduction

- ▶ R is a statistical programming language<sup>2</sup> / environment<sup>3</sup>
- ▶ Includes a plethora of statistical features and graphical tools
- ▶ Interpreted language: offers command line interpreter (no compilation necessary)
- ▶ Nevertheless often quite fast (partially implemented in C and Fortran)
- ▶ Easy to extend via packages.<sup>4</sup>
- ▶ Open Source under GNU GPL v2

<sup>2</sup> R is no general purpose language!

<sup>3</sup> According to R-manual: ““environment” is intended to characterize it as a fully planned and coherent system, rather than an incremental accretion of very specific and inflexible tools, as is frequently the case with other data analysis software.”

<sup>4</sup> Though R does a bad job at providing nice supporting tools. However, the community does via many useful packages.

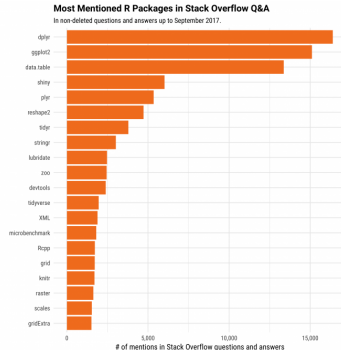


## C'mon! Yet another programming language?

### Why should I learn R? 2.0

- ▶ Since August 2021, ranked 14<sup>th</sup> in the TIOBE programming language popularity index.<sup>5</sup>
- ▶ Increasing popularity in industry
- ▶ Easy to learn, mostly fun to use the language (in particular for users with few or no programming experience)
- ▶ Few lines of code can actually realize a lot (see next slide)
- ▶ Always a good idea to know multiple programming language.<sup>6</sup>
- ▶ We will mostly apply existing implementations: no need for deep algorithmics (only if you want to; see advanced exercises in lecture notes 😊)

# The Impressive Growth of R (by Stack Overflow)

Study performed by Stack Overflow back in 2017 shows impressive growth of R:



Figures taken from Stack Overflow blog post<sup>7</sup>.

<sup>7</sup>


<https://stackoverflow.blog/2017/10/10/impressive-growth-r/>



## First impression

Achieve a lot with few lines of code:

```
> data(mtcars)
> X = mtcars[, c("qsec", "wt")]
> cl = kmeans(X, centers = 3, algorithm = "Lloyd")
> plot(X, pch = cl$cluster)
```



## Editors with R support

The interactive shell is nice, but useless for larger projects. We want an editor with ...

- ▶ ... nice features: syntax-highlighting, code-completion etc.,
- ▶ possibility to save code in files and
- ▶ organize/maintain our data analysis project(s)

## Nice editors

**R-Studio** powerful Integrated Development Environment (IDE) for R.

**Sublime-Text** Not an IDE! Very reduced, but minimalistic and lightning fast.

**Visual Studio Code** by Mircosoft; really good IDE with OKish R support.

## Installation

R does not ship with the common operating systems

**Windows** Download latest pre-compiled binary version from CRAN

**macOS** Download latest pre-compiled binary version or use Homebrew<sup>8</sup>:  
`brew upgrade`  
`brew install R`

**Linux** Depends on the flavour of you distribution, but your OS package manager should help. E.g., on Ubuntu  
`sudo apt update`  
`sudo apt install r-base`

**Some exotic OS** I have no idea 😊

---

<sup>8</sup> You might want to use Homebrew if you do not yet 😊

## Getting started: interacting with R

Start the interpreter by calling R from the command line<sup>9</sup> opens an interactive R session where you can type anything you want. At the *prompt* (`>`) we can perform calculations:

```
> 1 + 2 # this is a comment
## [1] 3

> 1 + 2 * 3 # R has operator precedence
## [1] 7

> # 1 + (2 * 3) # However, I like to make it explicit
> (1 + 2) * 3
## [1] 9

> log(10000) # function are called by passing arguments
## [1] 9.21034

> log(10000, base = 2)
## [1] 13.28771

> sqrt(25) # square root
## [1] 5

> 5^2 # power
## [1] 25
```

## Looking for help

- ▶ Call `help.start()` to enter the documentation entry point.<sup>10</sup>
- ▶ For a function, say `mean`, you can view the documentation via `?mean`.<sup>11</sup>
- ▶ Exact function name unknown: use `??mean`
- ▶ Looking for something similar via `apropos(...)`, e.g.,

```
> apropos("var")  
## [1] "all.vars"          "combine_vars"    "estVar"          "get_all_vars"  
## [5] "globalVariables"  "var"             "var.test"        "variable.names"  
## [9] "varimax"          "vars"
```

- ▶ Struggling with error message(s)?  
    ~> copy & paste into your favourite online search engine  
    (might help to add "r" to your search request as an additional keyword)
- ▶ Ask ChatGPT

---

<sup>10</sup> Unfortunately, the help page layout is from the past century.

<sup>11</sup> Often documentations are full of details and might be overwhelming on first sight

# Data types and classes

## Atomic data types / vector types

As in all other programming languages there are several data types:

`numeric` Real-valued numbers.

`integer` Integer numbers.

`logical` Boolean / logical values: TRUE and FALSE.

`character` Strings / concatenation of letters.

`factor` Special type of character which is internally stored as integer (for efficiency reasons); used for categorical variables.

`complex` Complex numbers.

`raw` Not discussed here.

## Vectors: first look

Vectors (remember the mathematical definition?) are elementary building blocks in R. The very basic function to create a vector is `c` for combine / concatenate:

```
> c(14.4, 18.4, 19.5, 25.4) # numeric
## [1] 14.4 18.4 19.5 25.4

> c(1, 2, 5, 6) # actually also numeric
## [1] 1 2 5 6

> c(1L, 2L, 5L, 6L) # integer
## [1] 1 2 5 6

> c("dog", "cat", "rat", "mouse") # character
## [1] "dog" "cat" "rat" "mouse"

> c(TRUE, FALSE, TRUE, FALSE, FALSE) # logical
## [1] TRUE FALSE TRUE FALSE FALSE

> x = 1:10 # save vector in variable
> x
## [1] 1 2 3 4 5 6 7 8 9 10

> x = c(x, 11) # append another element to
> x
## [1] 1 2 3 4 5 6 7 8 9 10 11
```



## Vectors: implicit coercion

All elements of a vectors are of the same type  $\leadsto$  in R all elements are coerced to the most general atomic data type:

```
> x = c(1, 2, 5, 6, 10.5)
> x
## [1] 1.0 2.0 5.0 6.0 10.5
```

```
> class(x)
## [1] "numeric"
```

```
> x = c(TRUE, "a", 5.43, 1)
> x
## [1] "TRUE" "a"      "5.43" "1"
```

```
> class(x)
## [1] "character"
```

```
> x = c(FALSE, 1)
> x
## [1] 0 1
```

```
> class(x)
## [1] "numeric"
```

## Vectors: explicit coercion

Sometimes it makes sense to explicitly convert/coerce a vector into another data format:

```
> x_char = c("19.7", "35.45", "29.34", "24.99", "23")
> class(x_char)
## [1] "character"

> x_num = as.numeric(x_char)
> x_num
## [1] 19.70 35.45 29.34 24.99 23.00

> class(x_num)
## [1] "numeric"

> x_int = as.integer(x_num) # information loss
> x_int
## [1] 19 35 29 24 23

> class(x)
## [1] "numeric"
```

## Vectors: sequencing

Generate sequences of numbers:

```
> seq(0, 1, by = 0.1)
## [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0

> seq(1, 0, by = -0.2)
## [1] 1.0 0.8 0.6 0.4 0.2 0.0

> seq(1, 5, length.out = 10)
## [1] 1.000000 1.444444 1.888889 2.333333 2.777778 3.222222 3.666667 4.111111
## [9] 4.555556 5.000000
```

## Vectors: repetitions

`rep(...)` for repetition is handy to create special types of vectors:

```
> rep(4, 10)
## [1] 4 4 4 4 4 4 4 4 4 4

> rep(4:6, 1:3)
## [1] 4 5 5 6 6 6

> rep(4:6, 3)
## [1] 4 5 6 4 5 6 4 5 6

> rep(4:6, each = 3)
## [1] 4 4 4 5 5 5 6 6 6

> # Use c to concatenate vectors
> x = rep(4, 3)
> c(10, x, rep(2:3, each = 3))
## [1] 10 4 4 4 2 2 2 3 3 3
```

## Working with vectors

R offers many predefined functions to work with vectors:

```
> x = rep(c(1, 2), each = 4); y = 1:8
> z = x + y; z # vector addition
## [1] 2 3 4 5 7 8 9 10

> 2 * x # scalar multiplication
## [1] 2 2 2 2 4 4 4 4

> x^2 # same as x * x
## [1] 1 1 1 1 4 4 4 4

> c(2, 3) * x # the shorter vector is repeated
## [1] 2 3 2 3 4 6 4 6

> sum(x) # sum :)
## [1] 12

> mean(x) # arithmetic mean
## [1] 1.5

> var(x) # sample variance
## [1] 0.2857143
```

## Numeric functions

Function	Description
<code>abs(x)</code>	Absolute value
<code>sqrt(x)</code>	Square root
<code>log(x)</code>	Natural logarithm (base $e$ )
<code>log10(x)</code>	Logarithm with base 10
<code>exp(x)</code>	Exponential function $e^x$
<code>cos(x), sin(x), ...</code>	Trigonometric functions
<code>ceiling(x)</code>	Round up: <code>ceiling(6.475)</code> is 7
<code>floor(x)</code>	Round down: <code>floor(6.489)</code> is 6
<code>trunc(x)</code>	Cut decimals: <code>trunc(2.99)</code> is 2
<code>round(x, digits=n)</code>	Regular rounding: <code>round(7.657, 2)</code> yields 7.67

## Statistical functions

Function	Description
<code>mean(x, na.rm=FALSE)</code>	Arithmetic mean of object <code>x</code>
<code>sd(x)</code>	Standard deviation of object <code>x</code>
<code>sd(x)</code>	Variance of object <code>x</code>
<code>mad(x)</code>	Median absolute deviation of values in <code>x</code>
<code>median(x)</code>	Median value of object <code>x</code>
<code>quantile(x, probs)</code>	Quantiles where <code>x</code> is the numeric vector whose quantiles are desired and <code>probs</code> is a numeric vector with probabilities in $[0, 1]$
<code>range(x)</code>	Range
<code>sum(x)</code>	Sum
<code>min(x)</code>	Minimum
<code>max(x)</code>	Maximum

Note: most of these functions have many more parameters!

## Exercises



1. Generate the following vectors in R using `seq`, `rep` and combinations thereof:

```
> c(4,5,6,4,5,6,4,5,6)
> c(1,1,1,1,2,2,2,2,3,3,3)
> c(0, 0.2, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0)
> c(1,1,2,2,3,3,8,9,8,9,8,9,8,9)
```

2. Consider the vector

```
> a = c(17.4, 19.4, 15.3, 19.3, 25.3, 16.4, 20.3, 18.6)
```

- ▶ Extract the 1<sup>st</sup>, 2<sup>nd</sup> and 4<sup>th</sup> elements and assign them to a new vector `b`
- ▶ Figure out how to calculate the length of a vector
- ▶ Extract the last 50% of the vector and assign it to a vector `c`



## Logical operations

Given a vector we can apply logical operations to compare vector(s) or a vector with a scalar:

```
> x = c(1, 6, 2, 8, 10, 4, 3); y = c(9, 4, 2, 5, 2, 18, 4)
>
> x < y # strictly less
## [1] TRUE FALSE FALSE FALSE FALSE TRUE TRUE

> x >= y # greater or equal
## [1] FALSE TRUE TRUE TRUE TRUE FALSE FALSE

> x == max(x) # equality
## [1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE

> !(x == max(x)) # negation
## [1] TRUE TRUE TRUE TRUE FALSE TRUE TRUE

> (x >= 3) & (x <= 5) # logical AND (component-wise)
## [1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE

> (y > 9) | (y < 3) # logical OR (component-wise)
## [1] FALSE FALSE TRUE FALSE TRUE TRUE FALSE
```

## Logical operations

**Attention:** for binary Boolean/logical operators there also exist "long" versions which only evaluate the first element!

```
> x = 1:4
> y = -1:2
> (x < 3) & (y < 2) # component-wise
## [1] TRUE TRUE FALSE FALSE

> (x < 3) && (y < 2) # only the first component is checked and no warning!

## Error in (x < 3) && (y < 2): 'length = 4' in coercion to 'logical(1)'

> all((x < 3) & (y < 2)) # are all component-wise comparisons true?
## [1] FALSE

> any(!((x < 3) & (y < 2))) # is at least one component-wise comparison false?
## [1] TRUE
```

Moreover, the longer version implements *short-circuit evaluation*, i.e., if we say `x && y` for two scalars `x` and `y`, and `x` evaluates to `TRUE`, `y` is not evaluated at all.<sup>12</sup>

<sup>12</sup> These semantics can be found in most languages.

## Missing data

R is build around data and unfortunately data is often missing ☹

Missing data is represented with the keyword NA in vectors<sup>13</sup>.

```
> age = c(24, 23, 21, 21, NA, 31, NA, 19)
> age
## [1] 24 23 21 21 NA 31 NA 19

> mean(age) # most functions return NA if there is at least one NA in the data
## [1] NA

> mean(age, na.rm = TRUE) # set na.rm = TRUE to skip NA values
## [1] 23.16667

> is.na(age) # # check if elements are NA (x == NA does not work!)
## [1] FALSE FALSE FALSE FALSE TRUE FALSE TRUE FALSE

> !is.na(age) # get only non-NA values
## [1] TRUE TRUE TRUE TRUE FALSE TRUE FALSE TRUE

> age[complete.cases(age)] # same
## [1] 24 23 21 21 31 19
```

---

<sup>13</sup> NA for Not Available.

## Vector subsetting by indexing

Access a subset of the vector elements by integer indices:

```
> x = c(10, 6, 3, 6, 3, 5, 9, 5, 15, 6)
> x[5] # single elements
## [1] 3

> x[-5] # all but the 5th element
## [1] 10 6 3 6 5 9 5 15 6

> x[-c(5, 7)] # all but the 5th and 7th element
## [1] 10 6 3 6 5 5 15 6

> x[c(1, 5, 10)] # multiple elements
## [1] 10 3 6

> x[1:2] = 20 # re-assign the first two elements
> x
## [1] 20 20 3 6 3 5 9 5 15 6
```

## Exercises



1. Consider the vector  $x = (49, 14, 25, 49, 14, 63, 65, 99, 56, 29)$ . Create this vector in  $\mathbb{R}$
2. Create a logical vector where the  $i^{\text{th}}$  entry is TRUE if the  $i^{\text{th}}$  entry of  $x$  is equal to the minimum value of  $x$
3. The function `rev` reverts the order of a vector. Use it to revert the first half of  $x$
4. Normalize/rescale the vector: i.e., from each element subtract the minimum and divide by the difference of maximum and minimum value (functions `min` and `max` will be useful)

## Vector subsetting by name

Vector elements can be named. The names can be used for subsetting:

```
> x = c("a" = 190, "b" = 20, "c" = 31)
> x[2] # by index of course still works
##    b
## 20
```

```
> x["a"] # single element
##    a
## 190
```

```
> ns = c("a", "b", rep("c", 3))
> x[ns]
##    a    b    c    c    c
## 190   20   31   31   31
```

```
> unname(ns) # drop names
## [1] "a" "b" "c" "c" "c"
```

```
> x["d"] # non-existing name -> NA
## <NA>
##    NA
```

## Vector subsetting with logical values

Remember the logical operations? We can also use logical vectors to index vectors.

```
> x = c(10, 6, 3, 6, 3, 5, 9, 5, 15, 6)
>
> large = x >= 6
> large
## [1] TRUE TRUE FALSE TRUE FALSE FALSE TRUE FALSE TRUE TRUE

> x[large]
## [1] 10 6 6 9 15 6

> x[x >= 6 & x <= 10]
## [1] 10 6 6 9 6

> x[rep(c(TRUE, FALSE), each = 5)]
## [1] 10 6 3 6 3
```

## How to obtain indices of items of interest

Occasionally it is useful to not get the actual components but their indices. Here, function `which(x)` shines:

```
> x = c(10, 6, 3, 6, 3, 5, 9, 5, 15, 6)
>
> x >= mean(x)
## [1] TRUE FALSE FALSE FALSE FALSE FALSE TRUE FALSE TRUE FALSE

> idx_large = which(x >= mean(x))
> idx_large
## [1] 1 7 9

> x[idx_large]
## [1] 10 9 15

> idx_min = which(x == min(x))
> idx_min
## [1] 3 5

> idx_min = which.min(x)
> idx_min
## [1] 3
```



## Some very useful functions

**Homework:** toy around with these functions. Check the documentation (`?function`).

Function	Description
<code>sort(x)</code>	Sort elements
<code>order(x)</code>	Indices of elements in sorted order
<code>unique(x)</code>	Vector of unique elements (removes duplicates)
<code>duplicated(x)</code>	Which elements of <code>x</code> are duplicates?
<code>which.min(x)</code>	Index of smallest element
<code>which.max(x)</code>	Index of largest element
<code>which(x)</code>	Indices of elements in <code>x</code> which are TRUE

## Exercises



1. Create a vector as follows:

```
> x = sample(1:5, size = 25, replace = TRUE)
```

2. Research what the function `sample` does.
3. Apply `duplicated(x)` and try to make sense out of it.
4. Create a vector which contains all unique elements of `x` in two different ways.
5. Generate a vector `y` which contains the indices of the duplicates of `x`.
6. Create a vector `y` which contains the indices of unique elements in the sorted version of `x`.

## Attributes

- ▶ Every object in R can have *attributes* (basically key-value pairs) attached. This is basically meta-data appended to the object.

```
> x = c(1:5, NA, NA, 6, 2, 6)
> attr(x, "len") = length(x)
> attr(x, "nna") = sum(is.na(x))
> x
## [1] 1 2 3 4 5 NA NA 6 2 6
## attr("len")
## [1] 10
## attr("nna")
## [1] 2

> attr(x, "nna") # read attribute
## [1] 2

> attributes(x) # read all attributes
## $len
## [1] 10
##
## $nna
## [1] 2

> attr(x, "len") = NULL # drop attribute
```

- ▶ Most attributes are dropped by operations (except for dim and/or names).

# Matrices

A matrix in R is like a matrix in mathematics:

```
> x = matrix(1:9, ncol = 3)
> x
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

> x[2, 3] # element in 2nd row and 3rd column
## [1] 8

> x[1, ] # first row (vector)
## [1] 1 4 7

> x[1, , drop = FALSE] # keep matrix structure even if we extract a single line
##      [,1] [,2] [,3]
## [1,]    1    4    7

> x[, 2] # first column (vector)
## [1] 4 5 6

> x[c(1, 3), 1:2] # sub-matrix
##      [,1] [,2]
## [1,]    1    4
## [2,]    3    6
```

# Matrices

Actually a matrix is just a vector with a dim attribute:

```
> attributes(x)
## $dim
## [1] 3 3

> y = 1:9
> attr(y, "dim") = c(3, 3) # "manually" transform to matrix
> y
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

> is.matrix(y)
## [1] TRUE
```

# Matrices

Since matrices are vectors we can apply all functions that work for vectors:

```
> x - 10
##      [,1] [,2] [,3]
## [1,]   -9   -6   -3
## [2,]   -8   -5   -2
## [3,]   -7   -4   -1

> mean(x)
## [1] 5

> x[x >= 5] # subsetting drops dim attribute
## [1] 5 6 7 8 9

> x[x %% 2 == 0] = 100 # assignment does not drop
> x
##      [,1] [,2] [,3]
## [1,]    1  100    7
## [2,]  100    5  100
## [3,]    3  100    9
```

## Functions on matrices

Function	Description
<code>A %% B</code>	Matrix product $A \cdot B$
<code>A * B</code>	Element-wise product of matrices $(A_{ij} \cdot B_{ij}, 1 \leq i, j, \leq n)$
<code>t(A)</code>	Matrix transposition
<code>eigen(A)</code>	Eigenvalues and Eigenvectors of $A$
<code>solve(A)</code>	(regular) inverse $A^{-1}$ of $A$
<code>solve(A, b)</code>	Solution $x$ of equation $Ax = b$
<code>rowMeans(A)</code>	Row-wise mean values
<code>colMeans(A)</code>	Column-wise mean values
<code>rowSums(A)</code>	Row-wise sums
<code>colSums(A)</code>	Column-wise sums
<code>nrow(A), ncol(A)</code>	Number of rows/columns of $A$
<code>dim(A)</code>	Dimension (i.e., rows and columns) of $A$
<code>diag(A)</code>	Vector of diagonal elements of $A$

# Lists

Lists are similar to vectors, but each element can have a different data type/class.

```
> l = list(name = "Jakob", lectures = c("DA1", "OR"), age = 37)
>
> l[[1]] # access element by index
## [1] "Jakob"

> l$name # access by name
## [1] "Jakob"

> l[["name"]] # alternative (good if name is stored in variable)
## [1] "Jakob"

> l[["name"]] = "Jakob B." # assignment works too
> l[["name"]]
## [1] "Jakob B."

> l[["lectures"]][1]
## [1] "DA1"
```

Often the return values of complex functions are named lists

~ multiple return values.



## Lists (cont.)

We can add elements to existing lists:

```
> l = list(name = "Jakob", lectures = c("DA1", "OR"), age = 37)
> l$surname = "Bossek"
> l
## $name
## [1] "Jakob"
##
## $lectures
## [1] "DA1" "OR"
##
## $age
## [1] 37
##
## $surname
## [1] "Bossek"
```

Concatenation works as well:

```
> l2 = c(1, list(degree = "PHD", studied = "computer science"))
> is.list(l2)
## [1] TRUE

> l2$degree
## [1] "PHD"
```

# Dataframes

Data frames are lists where every components has the same length:

```
> x = data.frame(  
+   id = 1:4,  
+   name = c("Max", "Sophie", "Jack", "Ted"),  
+   grade = c(5.0, 5.0, 4.0, 5.0))  
> x  
##   id  name grade  
## 1  1   Max     5  
## 2  2 Sophie    5  
## 3  3   Jack    4  
## 4  4    Ted    5  
  
> x$name  
## [1] "Max"    "Sophie" "Jack"   "Ted"  
  
> x[1:2, c("name", "id")] # equal to x[1:2, c(2, 1)]  
##      name id  
## 1     Max  1  
## 2 Sophie  2  
  
> x[1:2, ]$id # get id (alternatives: x[1:2, "id"] or x[1:2, 1])  
## [1] 1 2
```

# Dataframes: subsetting

```
> # ...  
> x$sex = c("M", "F", "M", "M") # add another variable  
> head(x, n = 3) # show first n rows only  
##   id  name grade sex  
## 1  1   Max     5   M  
## 2  2 Sophie    5   F  
## 3  3   Jack    4   M  
  
> tail(x, n = 1) # show last n rows only  
##   id name grade sex  
## 4  4   Ted     5   M  
  
> # Subsetting works the way we expect it to work  
> x[x$sex == "M", ] # get all males  
##   id name grade sex  
## 1  1   Max     5   M  
## 3  3   Jack    4   M  
## 4  4   Ted     5   M  
  
> x[(x$sex == "M") & (x$grade != 5), "name"] # get all males that failed the DA1 exam :(  
## [1] "Jack"
```

## Dataframes: subsetting

```
> # ...  
> x1 = x[(x$sex == "M") & (x$grade != 5), "name"] # Remember this?  
>  
> subset(x, sex == "M" & grade != 5, select = "name") # same result, nicer interface  
##   name  
## 3 Jack  
  
> # Some helpful functions (also work for matrices)  
> ncol(x) # number of columns  
## [1] 4  
  
> nrow(x) # number of rows  
## [1] 4  
  
> dim(x) # dimension (nr. of rows, nr. of columns)  
## [1] 4 4
```

## Dataframes: within

The `within` enables to evaluate an expression in an environment constructed from data.<sup>14</sup>

```
> data(airquality) # check docs with ?airquality
> head(airquality, n = 1)
##      Ozone Solar.R Wind Temp Month Day
## 1      41      190  7.4   67     5   1
```

We can modify the data as follows (hard to read)

```
> airquality$Month = factor(month.abb[airquality$Month])
> airquality$TempCel = round((airquality$Temp - 32) * 5/9, 1) # Fahrenheit to Celsius
> airquality$Day = NULL
```

### Alternative with `within`

```
> data(airquality)
> airquality = within(airquality, {
+   Month = factor(month.abb[Month])
+   TempCel = round((Temp - 32) * 5/9, 1) # Fahrenheit to Celsius
+   rm(Day)
+ })
```

---

<sup>14</sup> This is not limited to dataframes.

## Dataframes: import

Importing data from flat files in CSV-format:<sup>15</sup>

```
> # writing to a csv file using . as decimal point and ; as the "comma"
> data(mtcars)
> write.table(mtcars, file = "mtcars.csv", row.names = TRUE, sep = ";", dec = ".")
```

The content of `mtcars.csv`:

```
"mpg";"cyl";"disp";"hp";"drat";"wt";"qsec";"vs";"am";"gear";"carb"
"Mazda RX4";21;6;160;110;3.9;2.62;16.46;0;1;4;4
"Mazda RX4 Wag";21;6;160;110;3.9;2.875;17.02;0;1;4;4
```

```
...
> # now let us import the data
> mtcars = read.table("mtcars.csv", header = TRUE, sep = ";")
> head(mtcars , 2)
##           mpg  cyl  disp  hp  drat    wt  qsec vs  am gear carb
## Mazda RX4    21   6  160 110   3.9 2.620 16.46 0   1    4    4
## Mazda RX4 Wag 21   6  160 110   3.9 2.875 17.02 0   1    4    4
```

---

<sup>15</sup> CSV = Comma Separated Values. Very common flat-file format for rectangular data.

## Dataframes: combining

```
> names = c("Anton", "Jack", "Tobias", "Sophie")
> ages = c(21, 31, 26, 35)
>
> cbind(names, ages) # cbind = column bind
##      names      ages
## [1,] "Anton"    "21"
## [2,] "Jack"     "31"
## [3,] "Tobias"   "26"
## [4,] "Sophie"   "35"

> as.data.frame(cbind(names, ages))
##      names ages
## 1  Anton   21
## 2   Jack   31
## 3 Tobias   26
## 4 Sophie   35

> stud1 = list(name = "Anton", age = 21)
> stud2 = list(name = "Sophie", age = 35)
>
> rbind(stud1, stud2) # rbind = row bind
##      name      age
## stud1 "Anton"   21
## stud2 "Sophie"  35
```

## Data frames: aggregating

In data analysis we often split a data frame by some variables and calculate summary statistics, e.g., the mean, variance etc.

```
> data(mtcars)
>
> aggregate(
+   mtcars[, c("mpg", "cyl", "disp", "gear")], # data set
+   by = list(mtcars$cyl), # list of grouping variables
+   FUN = mean # aggregation function
+ )
##   Group.1      mpg cyl    disp  gear
## 1      4 26.66364   4 105.1364 4.090909
## 2      6 19.74286   6 183.3143 3.857143
## 3      8 15.10000   8 353.1000 3.285714
```



# Data frames: aggregating

## Splitting by multiple variables:

```
> aggregate(
+   mtcars[, c("mpg", "cyl", "disp", "gear")], # data set
+   by = list(mtcars$cyl, mtcars$gear), # grouping variable
+   FUN = function(x) {
+     c(mean = mean(x), sd = sd(x))
+   },
+   drop = TRUE # drop unused combinations of grouping values
+ )
##   Group.1 Group.2   mpg.mean    mpg.sd cyl.mean cyl.sd  disp.mean  disp.sd
## 1      4      3 21.5000000    NA      4      NA 120.100000    NA
## 2      6      3 19.7500000  2.3334524      6      0 241.500000  23.334524
## 3      8      3 15.0500000  2.7743959      8      0 357.616667  71.823494
## 4      4      4 26.9250000  4.8073604      4      0 102.625000  30.742699
## 5      6      4 19.7500000  1.5524175      6      0 163.800000   4.387862
## 6      4      5 28.2000000  3.1112698      4      0 107.700000  17.819091
## 7      6      5 19.7000000    NA      6     NA 145.000000    NA
## 8      8      5 15.4000000  0.5656854      8      0 326.000000  35.355339
##   gear.mean gear.sd
## 1      3      NA
## 2      3      0
## 3      3      0
## 4      4      0
## 5      4      0
## 6      5      0
## 7      5     NA
```

## Data frames: outlook

R ships with nice methods for subsetting (`subset`), aggregating (`aggregate`) etc., but data analysis become real fun with additional packages:

```
> suppressPackageStartupMessages(library(tidyverse))
> mtcars %>%
+   select(mpg, cyl, disp, gear) %>%
+   group_by(cyl, gear) %>%
+   summarize_all(list(mean = mean, sd = sd)) %>%
+   ungroup()
## # A tibble: 8 x 6
##       cyl  gear mpg_mean disp_mean mpg_sd disp_sd
##   <dbl> <dbl>   <dbl>    <dbl>  <dbl>  <dbl>
## 1     4     3    21.5      120.   NA     NA
## 2     4     4    26.9      103.   4.81   30.7
## 3     4     5    28.2      108.   3.11   17.8
## 4     6     3    19.8      242.   2.33   23.3
## 5     6     4    19.8      164.   1.55    4.39
## 6     6     5    19.7      145.   NA     NA
## 7     8     3    15.0      358.   2.77   71.8
## 8     8     5    15.4      326.   0.566  35.4
```

We will learn about the `tidyverse` suite next time!

## Exercises



1. Load the `mtcars` data set (`data(mtcars)`)
2. Subset all observations where the horsepower is among the highest 50% of the values
3. Subset all observations where the weight is in the interval  $[2, 3]$  and the number of cylinders is at least 6
4. Create a data frame that contains the first and last 5 observations of `mtcars`
5. Research how to create a random subset of observations (Hint: check `?sample`)
6. Add a new variable `wtcat` which takes two character values: "heavy" if the weight exceeds 4 and "light" otherwise.

## Factors aka categorical variables

Oftentimes variables have a limited number of categories, e.g., sex, field of study or, say, experience with R.

```
> sex_char = c("F", "M", "F", "F", "M")
> sex_char
## [1] "F" "M" "F" "F" "M"

> sex_fac = factor(sex_char) # convert to factor
> sex_fac
## [1] F M F F M
## Levels: F M

> sex_fac[2] = "F"
> sex_fac[2] = "other" # fails!

## Warning in `[<-,factor'('other', 2, value = "other")': invalid factor level, NA generated

> sex_char[2] = "other" # no problem at all
> sex_char
## [1] "F"      "other" "F"      "F"      "M"
```

## Factors: renaming categories

Sometimes it makes sense to rename factor levels:

```
> # no explicit levels
> sex_fac = factor(c("F", "M", "F", "F", "M"))
> sex_fac
## [1] F M F F M
## Levels: F M

> # level = category
> levels(sex_fac)
## [1] "F" "M"

> # we can also use it to rename
> levels(sex_fac) = c("Female", "Male", "Other")
> sex_fac
## [1] Female Male   Female Female Male
## Levels: Female Male Other



> sex_fac[2] = "Other"
> sex_fac
## [1] Female Other   Female Female Male
## Levels: Female Male Other
```

## Factors: example

Imagine we imported data where missing values were encoded differently (not the R way with NA):

```
> # A dash "-" indicates a missing value
> party = factor(c("SPD", "CDU", "CDU", "SPD", "-", "SPD", "Grüne"))
> # The dash is a factor level now
> party
## [1] SPD   CDU   CDU   SPD   -      SPD   Grüne
## Levels: - CDU Grüne SPD

> opar = par(mfrow = c(1, 2), cex.axis = 0.6) # some graphical parameter
> plot(party)
> levels(party)[1] = "Missing"
> plot(party)
```



## Exercises



1. Create the following vector where both a dash and an empty string represent missing values:

```
> x = c("a", "b", "c", "-", "", "c", "a", "b", "c", "-", "/")
```

2. Use logical expressions to replace all occurrences of "-" and "" with NA in x
3. Convert x to a factor
4. Plot a barplot. What do you observe? Is this behaviour always desirable?

## Ordered factors

Sometimes categories can be ordered, e.g., for *knowledge in R* we may have *good* > *medium* > *none*.

- ▶ We can tell R that a factor is indeed ordered.
- ▶ Used by plots, e.g., to arrange factor levels on axis respecting their order.<sup>16</sup>

```
> r_fac = factor(c("good", "medium", "medium", "none", "good"),
+   levels = c("good", "medium", "none"), ordered = TRUE)
> r_fac
## [1] good   medium medium none    good
## Levels: good < medium < none

> # comparison makes sense for ordered factors ...
> r_fac[1] < r_fac[2]
## [1] TRUE

> # ... but not for unordered factors
> sex_fac[1] > sex_fac[2]
## [1] NA
```

---

<sup>16</sup> Default is alphabetic order.



## Exercises



1. Load the mtcars dataset via `data(mtcars)`
2. Read the documentation of the data set (`?mtcars`)
3. Plot a histogram of the gross horsepower variable
4. Now we want to add a new variable `power` to the data. Use the function `cut` to convert gross horsepower to the categorical variable `power` with factor levels `(0, 150]`, `(150, 250]` and `(250, 350]` (see argument `breaks` of `cut`)
5. Plot a barplot of `power`
6. Rename the factor levels to *low*, *medium* and *high* and make the factor ordered

## Inspecting objects

Very helpful function `str` for structure gives a nice overview of the data types/classes of an object (applicable to all kinds of R objects):

```
> str(x)
## 'data.frame': 4 obs. of  4 variables:
##  $ id    : int  1 2 3 4
##  $ name   : chr   "Max" "Sophie" "Jack" "Ted"
##  $ grade : num   5 5 4 5
##  $ sex    : chr   "M" "F" "M" "M"

> str(1:4)
## int [1:4] 1 2 3 4

> str(factor(c("good", "bad", "bad", "bad"), levels = c("good", "bad"), ordered = TRUE))
## Ord.factor w/ 2 levels "good"<"bad": 1 2 2 2

> str(mean)
## function (x, ...)
```

Output

## Executing R scripts

If we have an R script, say `myscript.R`, with R code we can execute it in an interactive session by typing

```
> source("myscript.R")
```

We can also execute scripts from the command line. This is powerful if you want to, e.g.,

- ▶ Automate processes or
- ▶ Call R through other tools.

```
Rscript myscript.R
```

```
Rscript myscript.R > output.txt # redirect output to file (on unix only)
```

```
R -e 'install.packages("ggplot2", dependencies = TRUE)'
```

## Output

So far we just wrote the name of an object and it got printed to the console, e.g.:

```
> x = 10  
> x  
## [1] 10
```

This works perfectly fine in *interactive mode*, but not if we run a script via `Rscript myscript.R` or `source("myscript.R")`. In these cases cases like the 2nd line in the above listing will have no effect.

The following works as expected though:

```
> x = 10  
> print(x)  
## [1] 10
```

## Output: string formatting

Function `sprintf(fmt, ...)` expects a so-called *format-specifier*<sup>17</sup> and an arbitrary number of arguments that are parsed/used by the former.

```
> sprintf("Just a string.") # works perfectly fine
## [1] "Just a string."

> x = 10
> y = 19.45353325
> s = "DA1 is awesome!"
> sprintf("x has the value %i", x) # %i for integer
## [1] "x has the value 10"

> sprintf("x = %i, y = %f", x, y) # %f for fixed point decimal notation [-]mmm.ddd
## [1] "x = 10, y = 19.453533"

> sprintf("y = %.2f", y) # two decimal number after .
## [1] "y = 19.45"

> sprintf("Simple fact: %s", s) # %s for character
## [1] "Simple fact: DA1 is awesome!"
```

---


<sup>17</sup> Essentially a character with specific instructions embedded. The latter are substituted by the formatted arguments passed after the format string in order of appearance.

## Basic R graphics

## Basic R graphics: plot

R offers heaps of build-in plot functionality:


```
> data(mtcars)
> plot(mtcars$hp, mtcars$wt, xlab = "Gross horsepower", ylab = "Weight (in lbs)")
```





## Basic R graphics: barplot


```
> tab = table(mtcars$cyl)
> barplot(tab, cex.axis = 0.7, main = "Distribution of the nr. of cylinders")
```



## Basic R graphics: hist

For univariate numeric variables the histogram is a good starting point:


```
> hist(mtcars$displacement, xlab = "Displacement")
```



## Basic R graphics: hist

For univariate numeric variables the histogram is a good starting point:


```
> hist(mtcars$displ, breaks = 20, freq = TRUE, col = "grey67",  
+      xlim = c(0, 600), xlab = "Displacement")
```



## Basic R graphics: scatter-plot

For bivariate numeric observations:


```
> data(iris)
> plot(iris$Sepal.Length, iris$Petal.Length,
+       xlab = "Sepal length", ylab = "Petal length")
```



## Basic R graphics: scatter-plot

For bivariate numeric observations:


```
> cols = c("#8DB3D9", "#8D8DD9", "#B38DD9") # colors in RGB format
> plot(iris$Sepal.Length, iris$Petal.Length,
+   xlab = "Sepal length", ylab = "Petal length",
+   col = cols[iris$Species], pch = 14) # color by species (factor)
> grid(nx = 10, ny = 10, col = "lightgray", lty = "dotted", equilogs = TRUE)
> legend(7, 2.5, legend = c("setosa", "versicolor", "virginica"),
+   col = cols, pch = 14)
```



## Basic R graphics: matrix of plots

Function `par` serves to either query or set graphical parameters which are applied to all subsequently produced plots.

```
> # save old settings in opar
> opar = par(
+   mfrow = c(1, 3), # (1 x 3) matrix arrangement of plots
+   cex.axis = 0.9, # axis labels font size slightly reduced
+   oma = rep(0, 4) # size of outer margins
+ )
> hist(iris$Sepal.Length, xlab = "Sepal length")
> hist(iris$Petal.Width, xlab = "Petal width")
> barplot(table(iris$Species), xlab = "Species")
```



```
> par(opar) # reset settings
```

## Basic R graphics: saving plots

Plots can be exported in the most prominent formats: PDF<sup>18</sup>, PNG<sup>19</sup>, TIFF<sup>20</sup> etc.

```
> png("myfile.png", units = "cm", width = 8, height = 10, res = 72) # initialize output
> barplot(iris$Species)
> dev.off() # close device. I.e., write file.
>
> jpeg("myfile.jpeg", quality = 85) # quality adjusts JPG compression
> boxplot(iris$Sepal.Width ~ iris$Species)
> dev.off()
>
> # works also for collections
> pdf("myfile.pdf", width = 700, height = 400)
> opar = par(
+   mfrow = c(1, 2), # (1 x 2) matrix arrangement of plots
+ )
> hist(iris$Sepal.Length, xlab = "Sepal length")
> hist(iris$Petal.Width, xlab = "Petal width")
> par(opar) # reset settings
> dev.off()
```

---

<sup>18</sup> PDF = Portable Data Format; vector-based format.

<sup>19</sup> PNG = Portable Network Graphics; pixel-based format.

<sup>20</sup> TIFF = Tag Image File Format.

## Exercises



1. Install and load the ggplot2 package with

```
> install.packages("ggplot2", dep = TRUE)
> library(ggplot2)
```

2. Load the diamonds dataset via `data(diamonds)`

3. Take a random sample of 500 diamonds:<sup>21</sup>

```
> x = diamonds[sample(1:nrow(diamonds), size = 500), ]
```

4. Toy around with the learned visualizations. Some ideas:

- ▶ Generate histograms of price and or carat side by side.
- ▶ Draw scatter-plots of x vs. y or price vs. carat. In the latter plot color the points by cut and add a legend; figure out how to omit the border of the legend.

---

<sup>21</sup> The data set contains 53 940 observations. It may take some time to render plots for that many observations. Try it!



## Basic R graphics: conclusion

- ▶ Simple, yet powerful.
- ▶ We just scratched the surface here!  
    ~→ Essentially we can achieve anything with base R graphics.
- ▶ However, some actions are not directly supported:
  - ▶ Facetting (multiple panels).
  - ▶ Error bars.
  - ▶ Legends need to be crafted manually.
  - ▶ etc.
  - ▶ Pro: generation of base R graphics is fast.
- ▶ Later we will learn about `ggplot` (Wickham 2009), a sophisticated visualization library/package.<sup>22</sup>

---

<sup>22</sup> It is not the holy grail, but way better and visually appealing.

Control flow

## Conditional statements

Conditions are elementary building blocks of (procedural) programming: do something if certain logical conditions hold and do something else if they do not hold.

```
> x = c(1, 2, 3, 3, 2, 1, 5, 3)
> (x %% 2) # modulo operator -> rest of integer division with 2
## [1] 1 0 1 1 0 1 1 1

> (x %% 2) == 0 # ith entry is TRUE if ith number is even
## [1] FALSE TRUE FALSE FALSE TRUE FALSE FALSE FALSE

> all((x %% 2) == 0) # TRUE if all elements of logical vector are TRUE
## [1] FALSE

> if (all((x %% 2) == 0)) {
+   "all numbers even"
+ } else {
+   "at least one number is odd"
+ }
## [1] "at least one number is odd"
```

## Conditional statements

If, else if, . . . , else if, else:

```
> x = c(43.3, 15.3, NA, NA, 34.3, Inf, 19.3, 15.4)
>
> xna = is.na(x)
> xinf = is.infinite(x)
>
> if (any(xna) & any(xinf)) {
+   # print "prints" its argument to the console even if the script is not executed
+   # interactively, but is loaded from an R-script
+   print("There are NAs and infinite values")
+ } else if ((sum(xna) > 1) & (all(x < 40))) {
+   print("There is at least one NA and all values re at most 40")
+ } else {
+   print("Nothing interesting to report")
+ }
## [1] "There are NAs and infinite values"
```

## Conditional statements: invalid inputs

The condition needs to be or evaluate to a single Boolean value!

```
> if ("char") 1

## Error in if ("char") 1: argument is not interpretable as logical

> if (c()) 1

## Error in if (c()) 1: argument is of length zero

> if (mean(c(1, NA, 4, 3) > 2)) 1

## Error in if (mean(c(1, NA, 4, 3) > 2)) 1: argument is not interpretable as logical

> if (c(TRUE, FALSE, TRUE)) 1

## Error in if (c(TRUE, FALSE, TRUE)) 1: the condition has length > 1
```

Since R 3.5.0. we can turn the latter to throw an error via

```
> Sys.setenv("_R_CHECK_LENGTH_1_CONDITION_" = "true") # defaults to "false"
> if (c(TRUE, FALSE)) 1

## Error in if (c(TRUE, FALSE)) 1: the condition has length > 1
```

## Conditional statements: vectorized if

The last example on the previous slide is actually not constructed at all: imagine we want to modify each element of a vector if some condition on this very element holds:

```
> x = c(6, 2, 6, 3, 2, 6, 2, 3, 8, 9, 1)
>
> # categorize into odd/even numbers (the laborious way)
> y = x
> for (i in seq_along(y)) {
+   if (x[i] %% 2 == 0) {
+     y[i] = "even"
+   } else {
+     y[i] = "odd"
+   }
+ }
> y
## [1] "even" "even" "even" "odd"  "even" "even" "even" "odd"  "even" "odd"
## [11] "odd"

> # and the R way
> y = ifelse(x %% 2 == 0, "even", "odd")
```

## For loop

Repetitive tasks are tedious. Programming languages like them a lot! 😊

```
> for (i in 1:3) {  
+   print(i^2)  
+ }  
## [1] 1  
## [1] 4  
## [1] 9  
  
> l = list(a = c(1, 3, 4), b = c(4, 5, NA), c = c(7,3,6))  
> names(l) # returns the names  
## [1] "a" "b" "c"  
  
> means = list()  
> for (name in names(l)) {  
+   means[[name]] = mean(l[[name]], na.rm = TRUE)  
+ }  
> unlist(means) # "flatten" the list  
##      a      b      c  
## 2.666667 4.500000 5.333333
```

## For loop

It is usually better to use `seq_along` to avoid hard to find errors if the sequence we iterate over is empty:

```
> x = c(2, 5, 2, 3, 3, 5)
> x = x[x > 10] # empty vector
>
> # Bad
> for (i in 1:length(x)) { # length(x) is 0
+   paste0(i, ": ", x[i]) # concatenate elements to character
+ }
>
> # Good: use seq_along(...)
> for (i in seq_along(x)) { # length(x) is 0
+   paste0(i, ": ", x[i]) # concatenate elements to character
+ }
>
> # for illustration
> 1:c() # not what we would like to have in a for loop

## Error in 1:c(): argument of length 0

> seq_along(c()) # this look better
## integer(0)
```



## For loop: next and break

We can skip iterations or break the entire loop earlier if some condition is satisfied:

```
> for (i in c(1, 2, NA, 4:10)) {  
+   if (is.na(i))  
+     next  
  
+   print(i)  
  
+   if (i >= 8)  
+     break  
+ }  
## [1] 1  
## [1] 2  
## [1] 4  
## [1] 5  
## [1] 6  
## [1] 7  
## [1] 8
```

## While loop

Repeat something as long as a logical condition is (not) met:

```
> # In this situation a for loop would be the better choice
> x = matrix(runif(100), nrow = 5)
> i = 1
>
> while (i <= nrow(x)) {
+   mean(x[i, ])
+   i = i + 1 # need to increase manually. Otherwise: infinite loop
+ }
>
> x = runif(1)
> while(x > 0.1) { # sample random number in [0,1] until a number < 0.1 is sampled
+   x = runif(1)
+ }
```

## Loops are slow in R, aren't they?

**Oftentimes you will read:** avoid loops in R. They are slow as hell!

### Golden rules to not slow down your R code with loops

- ▶ Use R's *vectorization*. I.e., if there exists a build-in vectorized function, use it!
- ▶ Avoid using `c` or `cbind/rbind` in loops to grow objects  
     $\leadsto$  R needs to re-allocate memory in every iteration.<sup>23</sup>
- ▶ Instead, try to pre-allocate memory in advance.

### Still a loop is unavoidable?

- ▶ Allocate memory in advance, e.g. `x = numeric(1000)` if you know that you will populate 1 000 positions of a numeric vector.
- ▶ If it is really a bottleneck  $\leadsto$  implement function in C or C++.

---

<sup>23</sup> If allocation takes linear time, a linear number of re-allocations takes quadratic time.

## R loops vs. build-in functions: study I

**Study:** Given a length  $n$  vector  $x$  let us build a vector  $y$  which contains the *cumulative sum*,<sup>24</sup> i.e.,

$$y_i = \sum_{j=1}^i x_j, i = 1, \dots, n.$$

```
> x = rnorm(100000) # data generation: N(0, 1), i.e., normally distributed randoms
> n = length(x)
>
> system.time(cumsum(x)) # build-in C-based function (ultra fast)
##      user      system elapsed
##       0         0         0

> system.time({ # BAD: growing vector in each iteration (ultra slow)
+   cs = c(x[1])
+   for (i in 2:n) {
+     cs = c(cs, cs[i - 1] + x[i])
+   }
+ })
##      user      system elapsed
##    7.101     1.179     8.297
```

---

<sup>24</sup> Note that calculations are not independent here, so vectorization is not an option.

## R loops vs. build-in functions: study I

Here are better loop solutions using pre-allocated memory:

```
> system.time({ # GOOD: pre-allocate memory in advance (good performance)
+   cs = numeric(n)
+   cs[1] = x[1]
+   for (i in 2:n) {
+     cs[i] = cs[i - 1] + x[i]
+   }
+ })
##      user  system elapsed
##    0.005   0.000   0.005
```

```
> system.time({
+   cs = numeric(n)
+   cs[1] = x[1]
+   i = 2
+   while (i <= n) {
+     cs[i] = cs[i - 1] + x[i]
+     i = i + 1
+   }
+ })
##      user  system elapsed
##    0.007   0.000   0.008
```

## R loops vs. build-in functions: study II

Now let us check the performance of vectorizable operations. For a given vector  $x$  we want to calculate a simple linear transformation  $a \cdot x + b$  for some constants  $a, b \in \mathbb{R}$ :

```
> x = runif(1000000) # uniform random numbers in [0, 1]
> n = length(x)
> a = 10
> b = 100
>
> system.time({y = a * x + b}) # build-in C-based function (ultra fast)
##      user  system elapsed
##    0.001    0.001    0.001

> system.time({y = sapply(x, function(e) a * e + b)}) # vectorized (slower than loop!)
##      user  system elapsed
##    0.490    0.021    0.512

> system.time({ # Not bad at all
+   y = x
+   for (i in 1:n) {
+     y[i] = a * y[i] + b
+   }
+ })
##      user  system elapsed
##    0.038    0.001    0.039
```

# Functions

## Reusable code

Functions are the most basic building blocks if it comes to reusable code. For those who are not familiar with the concept: imagine we want to calculate  $\sum_{i=1}^n x_i \cdot y_i$  over and over again for different  $x$  and  $y$ .

```
> x = 1:10; y = 11:20
> s = 0
> # this is actually a bad example for a loop, but helps to state the problem here
> for (i in 1:10) {
+   s = s + (x[i] * y[i])
+ }
>
> x = c(10, 24, 53); y = c(29, 10, 4)
> s = 0
> for (i in 1:10) {
+   s = s + (x[i] * y[i])
+ }
>
> # and so on ...
```

Pretty tedious, error-prone and redundant isn't it?



# Functions

Input arguments  $\longrightarrow$  **Function**  $\longrightarrow$  Return value(s)

- ▶ Functions encapsulate code that solves some interesting sub-task and is general enough to be reused.
- ▶ Allow for automation of common tasks.
- ▶ It is much like a math. function. E.g.

$$f : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R} \text{ with } f(x, y) \mapsto \sum_{i=1}^n x_i \cdot y_i.$$

- ▶ The good thing is: once we have a useful function, we no longer care about its exact implementation, but only about the function *interface/signature*. I.e., what does the function expect as input and what does it return as output?
- ▶ R functions are objects!
  - $\leadsto$  Functions can be assigned to variables, stored in vectors, passed down to other functions, etc.

## Functions: definition

Let's write a function for our problem:  $\sum_{i=1}^n x_i \cdot y_i$ .

```
> sumprod = function(x, y) {  
+   r = sum(x * y)  
+   return(r)  
+ }  
>  
> sumprod(c(29, 10, 4), c(10, 24, 53))  
## [1] 742  
  
> sumprod(x = 1:10, y = 11:20)  
## [1] 935  
  
> sumprod(x = c(10, 24, 53), c(29, 10, 4))  
## [1] 742  
  
> sumprod(y = c(29, 10, 4), x = c(10, 24, 53))  
## [1] 742
```

## Functions: scoping (simplified)

**Scoping** = process of finding the value associated with a name.

In functions we can (and often do<sup>25</sup>) redefine objects / variables names; the inner-most definition is relevant.

```
> x = 1; y = 10
> f = function() {
+   x = 100 # this variable "masks" the definition outside the function
+   c(x, y)
+ }
> f()
## [1] 100 10
```

```
> x
## [1] 1
```

```
> f = function() {
+   x = x + 1
+   x
+ }
> f()
## [1] 2
```

```
> x
## [1] 1
```

## Lazy evaluation

*Lazy evaluation* means that an expression is evaluated only if the expression (calculation, value of object) is actually used.

```
> lazyfun = function(x, y) {  
+   if (x < 10)  
+     x + 10  
+   else  
+     x + y[1]  
+ }  
>  
> lazyfun(10) # works since y is never used and thus not evaluated in lazyfun  
  
## Error in lazyfun(10): argument "y" is missing, with no default
```

Powerful tool! Helps to avoid costly calculations if the result won't be used.

```
> lazyfun(4)  
## [1] 14  
  
> lazyfun(10) # y not passed, but function tries to use it  
  
## Error in lazyfun(10): argument "y" is missing, with no default
```

## Lazy evaluation

The concept of lazy evaluation is a powerful tool! It helps to avoid costly calculations if the result(s) won't be used.

```
> lazyfun = function(x, y) {  
+   if (x < 10)  
+     x + 10  
+   else  
+     x + y # y is only touched if x is smaller 10  
+ }  
>  
> lazyfun(4) # works since y is never used and thus not evaluated in lazyfun  
## [1] 14  
  
> lazyfun(20) # errors, since y not passed, but function tries to use it  
  
## Error in lazyfun(20): argument "y" is missing, with no default  
  
> system.time({lazyfun(4, var(runif(10000000)))}) # costly evaluation skipped  
##      user system elapsed  
##      0      0      0  
  
> system.time({lazyfun(20, var(runif(10000000)))}) # costly evaluation not skipped  
##      user system elapsed  
## 0.799 0.092 0.894
```

## Lazy evaluation

Lazy evaluation is not limited to functions! It is a language concept.

```
> if (TRUE) {  
+   print(100)  
+ } else {  
+   print(z) # z not defined, but this branch is never executed  
+ }  
## [1] 100  
  
> x = 1:1000000  
> if (x[1] == 1 || (sqrt(sum(x^2)) > 500)) { # costly second condition skipped  
+   print("Second condition not checked.")  
+ }  
## [1] "Second condition not checked."
```

## Functions: invocation

The usual way to call a function is by using the familiar notation. E.g.:

```
> x = runif(10) # some random numbers
> mean(x, na.rm = TRUE, trim = 0.1)
## [1] 0.6346447
```

Occasionally it happens that our arguments are already stored in a list:

```
> args = list(x, na.rm = TRUE, trim = 0.1)
> do.call(mean, args) # call mean fun with parameter given in args
## [1] 0.6346447
```

Helpful to glue together dataframes:

```
> x = data.frame(x = 1, y = "a")
> y = data.frame(x = 2, y = "b")
> do.call(rbind, list(x, y))
##   x y
## 1 1 a
## 2 2 b
```

## Functions: composition

Often we apply a sequence of functions to data:

1. **Nesting**: useful for short sequences.

```
> x = runif(10)
> y = runif(10)
> sqrt(sum((x - y)^2))
## [1] 1.360139
```

2. **Intermediate objects**: can become tedious.

```
> tmp = (x - y)^2
> tmp = sum(tmp)
> sqrt(tmp)
## [1] 1.360139
```

3. **Piping**:<sup>26</sup> using the magnificent magrittr package.<sup>27</sup>

```
> library(magrittr)
> (x-y)^2 %>%
+   sum() %>%
+   sqrt()
## [1] 1.360139
```

---

<sup>26</sup> Similar to pipe operator in unix command line.

<sup>27</sup> We will make extensive use of piping later.



## Exercises



1. Write a function `means(x)` which expects a numeric matrix `x` and returns a vector where the  $i^{\text{th}}$  entry corresponds to the arithmetic mean of the  $i^{\text{th}}$  column of `x`.
2. Modify `means` such that it expects another argument `of`. If `of="column"` the function shall behave as in (1). If `of="row"` the row-wise means shall be calculated instead.
3. The Fibonacci-numbers<sup>28</sup> is a famous recursive infinite sequence of numbers following the building rule

$$F_1 = 0, F_2 = 1 \text{ and } F_n = F_{n-1} + F_{n-2} \text{ for } n \geq 3.$$

Write a function `fibonacci(n)` which returns a vector of the first  $n$  Fibonacci numbers.

---

<sup>28</sup> [https://en.wikipedia.org/wiki/Fibonacci\\_number](https://en.wikipedia.org/wiki/Fibonacci_number)

## The \*apply-function family

Often we have an  $n$ -dim. object  $x = (x_1, x_2, \dots, x_n) \in \mathcal{D}^n$  and a function  $f : \mathcal{D} \rightarrow \mathcal{E}$  that we want to apply to each element of  $x$ :

$$f(x) = f((x_1, x_2, \dots, x_n)) \mapsto (f(x_1), f(x_2), \dots, f(x_n)).$$

E.g.

- ▶ Apply a function to each element of a vector.
- ▶ Apply a function to each column (or row) of a matrix.
- ▶ Apply a function to each matrix in a list.
- ▶ etc.

## The \*apply-function family

`lapply(data, function)` is the most generic: given a vector/list data it applies function function to every component and returns a list or vector:

```
> x = list(a = 1:10, b = 10:20)
>
> # in order to calculate the the maximum of each component we could do
> x_max = numeric(2)
> for (i in 1:2) {
+   x_max[i] = max(x[[i]])
+ }
> x_max
## [1] 10 20
```

```
> lapply(x, max) # readable alternative
## $a
## [1] 10
##
## $b
## [1] 20
```

```
> # here we pass an "anonymous" function to calculate the range
> lapply(x, function(e) max(e) - min(e))
## $a
## [1] 9
##
## $b
```

## The \*apply-function family

`apply(data, MARGIN, function)` works on matrices. The second parameter `MARGIN` indicates if the function should be applied row- or colwise:

```
> x = matrix(1:9, ncol = 3)
> x
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

> apply(x, 1, sum) # rowwise sum
## [1] 12 15 18

> apply(x, 2, var) # colwise variance
## [1] 1 1 1
```

## Code style

Actually, the R interpreter does not care whether your code looks like this ...

```
> x=c(24, 23, 21, 21, NA, 31, NA, 19)
> f2=function(x, na.rm=FALSE){
+   if(na.rm){x = x[complete.cases(x)]}
+   return(sum(x)/length(x))}
> f(x)
```

... Or ...

```
> age = c(24, 23, 21, 21, NA, 31, NA, 19)
> average = function(x, na.rm = FALSE) {
+   if(na.rm) { # drop NAs if argument is TRUE
+     x = x[complete.cases(x)]
+   }
+   return(sum(x) / length(x))
+ }
> average(age)
```

But you should care! ☺

↪ Use e.g., Google's R Style Guide<sup>29</sup> for your own good!

<sup>29</sup> <https://google.github.io/styleguide/Rguide.html>

## Some Dos and Don'ts

- ▶ Use a style guide.
- ▶ Use speaking names for objects<sup>30</sup>:

```
> x = c(24, 25, 23, 31, 54, 31) # bad
> ages = c(24, 25, 23, 31, 54, 31) # good
> f = function(x) { ... } # bad
> rotate_vector = function(x, theta) { ... } # good
```

- ▶ Avoid overwriting existing objects, e.g., by defining a variable `mean` or a function `c`.
- ▶ Document your code:

```
> k = 3 # initialize variable k to 3 (BAAAD!)
> k = 3 # number of clusters for k-means algorithm (good)
>
> # ToDo: add documentation (BAAAD! :))
> rotate_vector = function(x, theta) { ... }
>
> # Rotates a numeric vector counterclockwise by a given angle.
> # @param x [numeric] Numeric input vector of length at least 2.
> # @param theta [numeric] Rotation angle in degrees.
> # @return [numeric] Rotated vector x. (GOOD!)
> rotate_vector = function(x, theta) { ... }
```

30

I admit we did not do it here very often. We used rather short names to save precious space.

## Packages

... are collections<sup>31</sup> of functions and data.

Packages can be installed from the Comprehensive R Archive Network (CRAN) via

```
> install.packages("devtools", dep = TRUE)
```

Development versions, e.g., from Git(Hub) repositories can be installed via

```
> devtools::install_github("jakobbossek/smoof", dep = TRUE)
```


---

<sup>31</sup> Basically software libraries.

## Packages (cont.)

Once a package is installed it can be loaded and all exported functions and datasets can be used:

```
> library(smoof)
> # the prefix 'smoof::' is not necessary, but useful to keep track where the function
> # is implemented
> fn = smoof::makeAckleyFunction(dimensions = 2)
> plot(fn, show.optimum = TRUE, render.levels = FALSE)
```





Advanced: Parallelization

## Parallelization in R: The Sad News

One of R's drawbacks: R will use one core regardless of how many cores the CPU exposes.

### How to deal with it?

- ▶ Simply use another programming language, e. g., Java, C, ...
- ▶ Interface R from C or Python.
- ▶ Compile the 64Bit Version using special build options.
- ▶ Use nice packages developed by the R community (see CRAN HPC task View).

### Types of parallelism

**Implicit parallelism** parallelism hidden from user, system abstracts it away.

**Explicit parallelism** user need to handle parallelism explicitly by using special directives.

# Implicit Parallelism

## Implicit parallelism?

Parallelism is hidden from the user, no user-requests/directives needed. Thus parallelism is automatically exploited.

- ▶ Paralellized computation of mathematical functions in Blas library.
- ▶ Installation on windows systems is simple: replace Rblas.dll with compiled one.
- ▶ Installation on unix systems: configure R with `-with-blas` flag and compile afterwards; fastest implementation is OpenBlas (hand optimized assembler code).

# Explicit Parallelism

## Explicit parallelism?

Parallelism is activated by the user calling specific directives.

## Drawbacks and pitfalls

- ▶ User needs to care for distribution of objects, loading of packages in the slave jobs.
- ▶ Logging is difficult.
- ▶ Identifying errors is even harder than in sequential mode.
- ▶ Random Number Generators (RNGs) must be initialized occasionally (this is of utmost importance in statistical computing).

## Explicit Parallelism: `parallel`

First simple package:

- ▶ Package shipped with R since version 2.14.0.
- ▶ Slight extension of `snow` and `multicore`.
- ▶ Includes parallel Random Numbers Generators (RNG).
- ▶ Main entry point is parallel version of `lapply`.
- ▶ Single CPU usage (`multicore`) or several machines (`snow`).
- ▶ On windows only in socket mode ☹.

## Explicit Parallelism: parallel (cont.)

```
> library(parallel)
> n.cores = 2 # number of cores
>
> x = 1:100
> x = split(x, rep (1:n.cores, each = length(x) / n.cores))
> str(x)
## List of 2
## $ 1: int [1:50] 1 2 3 4 5 6 7 8 9 10 ...
## $ 2: int [1:50] 51 52 53 54 55 56 57 58 59 60 ...

> y = unlist(mclapply(x, sum)) # parallel sum
> print(y)
##      1      2
## 1275 3775

> sum(y)
## [1] 5050
```

## Explicit Parallelism: parallel (cont.)

Example:  $k$ -means clustering on mtcars data.<sup>32</sup>

```
> f = function(i) {  
+   suppressPackageStartupMessages(library(mlr)) # load mlr package  
+   lrn = makeLearner("cluster.kmeans", centers = 2) # define k-means algorithm  
+   data(mtcars, package = "datasets")  
+   cluster.task = makeClusterTask(data = mtcars)  
+   mod = train(lrn, task = cluster.task) # calculate clustering  
+ }  
> n.cores = detectCores()  
> system.time({lapply(1:n.cores, f)}) # not parallel  
  
## Error in requirePackages(package, why = stri_paste("learner", id, sep = " "), : For learner cluster.kmeans please install the  
## following packages: clue  
## Timing stopped at: 0.319 0.032 0.955  
  
> system.time({mclapply(1:n.cores, f)}) # parallel  
  
## Warning in mclapply(1:n.cores, f): all scheduled cores encountered errors in user code  
## user system elapsed  
## 0.001 0.023 0.023
```

<sup>32</sup> The difference in runtime is almost negligible in this example since mtcars is a very small data set. However, if on larger data one run of  $k$ -means takes a minute you will experience a massive, close to optimal speed-up if you perform at most as many runs as there are CPU cores.

## Explicit parallelism: `parallelMap`

- ▶ Major drawback of methods seen so far: backend change requires change of code in general ☹
- ▶ R package `parallelMap`<sup>33</sup> is another wrapper
- ▶ Single function to learn: `parallelMap`
- ▶ No need to change code if backend changes
- ▶ Configurable via options
- ▶ Support for most important parallelization modes, i. e., multicore machines, socket mode, MPI and HPC

---

<sup>33</sup> <https://github.com/berndbischl/parallelMap>



## Explicit parallelism: parallelMap (cont.)

```
> library(parallelMap)
> # start in socket mode
> parallelStartSocket(cpus = 2L)

## Starting parallelization in mode=socket with cpus=2.

> f = function(i) {
+   res = summary(runif(1e6))
+ }
> parallelMap(f, 1:2) # like lapply(1:2, f)

## Mapping in parallel: mode = socket; level = NA; cpus = 2; elements = 2.
## [[1]]
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## 0.0000001 0.2502242 0.4998253 0.4997979 0.7497259 0.9999988
##
## [[2]]
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## 0.0000022 0.2496570 0.4994759 0.4995947 0.7492900 0.9999994

> parallelStop()

## Stopped parallelization. All cleaned up.
```

## Topics not covered

These slides just scratch the surface:

- ▶ The many peculiarities of R.<sup>34</sup>
- ▶ Object-oriented programming (S3, S4, R6).
- ▶ Metaprogramming.
- ▶ Advanced functional programming.
- ▶ Literate programming for writing reports and presentations where text, code and evaluation are intermingled.
- ▶ Interfacing C natively or C++ via Rcpp.
- ▶ Package development.

Work through Advanced R by H. Wickham (Wickham 2014) if you are interested in more details.

---

<sup>34</sup> IMHO it is not useful to teach this. These are things that come with experience.

# Literature Recommendations

## Literature Recommendations

1. Gareth James, Daniela Witten, Trevor John Hastie & Robert Tibshirani (2021). An Introduction to Statistical Learning. 2nd Edition. Springer. (James et al. 2013)★  
~> Brief R intro; R labs at the end of each chapter.
2. Hadley Wickham & Garrett Grolemund (2017). R for Data Science: Import, Tidy, Transform, Visualize, and Model Data. O'Reilly. (Wickham and Grolemund 2017)
3. Hadley Wickham (2016). ggplot2: Elegant Graphics for Data Analysis. Springer. (Wickham 2009)
4. Hadley Wickham (2014). Advanced R. Chapman & Hall/CRC The R Series. Taylor & Francis. (Wickham 2014)
5. Hadley Wickham (2015). R Packages. 1st Edition. O'Reilly Media. (Wickham 2015)

## What comes next?

Overview of the upcoming week:

Date	Content	Speaker
21 Oct.	Brief introduction to R ✓	Jakob
22 Oct.	Tutorial on math foundations	Jakob
28 Oct.	Preprocessing of data (with R)	Jakob
29 Oct.	Tutorial on R	Raphael
⋮	⋮	⋮

Wrap-Up

# Wrap-Up

## Today's content

R introduction (basic concepts, loops, functions etc.)

## Your task(s)

- ▶ Download and install both R<sup>35</sup> and R Studio<sup>36</sup>.
- ▶ Work through the R-intro in James et al. 2013.
- ▶ Work through the presentation slides and do the many exercises.
- ▶ Work on exercise sheet 02.

---

<sup>35</sup> URL: <https://www.r-project.org>

<sup>36</sup> URL: <https://www.rstudio.com>

## References I

- Wickham, Hadley (2009). *ggplot2: elegant graphics for data analysis*. Springer New York. ISBN: 978-0-387-98140-6.
- (2014). *Advanced R*. Chapman & Hall/CRC The R Series. Taylor & Francis. ISBN: 9781466586963. URL: <https://books.google.de/books?id=PFHFNAEACAAJ>.
- James, Gareth et al. (2013). *An Introduction to Statistical Learning: with Applications in R*. Springer.
- Wickham, Hadley and Garrett Golemund (Jan. 2017). *R for Data Science: Import, Tidy, Transform, Visualize, and Model Data*. 1st ed. O'Reilly Media. ISBN: 1491910399.
- Wickham, Hadley (2015). *R Packages*. 1st. O'Reilly Media, Inc. ISBN: 1491910593.